

```
# Licensed under the Apache License, Version 2.0 (the
# "License"); you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
# https://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
```

 Open in Colab

```
import csv
import tensorflow as tf
import numpy as np
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
!wget --no-check-certificate \
  https://storage.googleapis.com/laurencemoroney-blog.appspot.com/bbc-text.csv \
  -O /tmp/bbc-text.csv
```

```
--2020-09-19 19:42:51-- https://storage.googleapis.com/laurencemoroney-blog.appspot.com
Resolving storage.googleapis.com (storage.googleapis.com)... 173.194.214.128, 173.194.214.128
Connecting to storage.googleapis.com (storage.googleapis.com)|173.194.214.128|:443... cc
HTTP request sent, awaiting response... 200 OK
Length: 5057493 (4.8M) [application/octet-stream]
Saving to: '/tmp/bbc-text.csv'
```

```
/tmp/bbc-text.csv 100%[=====>] 4.82M --.-KB/s in 0.05s
```

```
2020-09-19 19:42:52 (105 MB/s) - '/tmp/bbc-text.csv' saved [5057493/5057493]
```

```
vocab_size = 10000
embedding_dim = 16
max_length = 120
trunc_type = "post"
padding_type = "post"
oov_tok = "<OOV>"
training_portion = .8
```

```
sentences = []
labels = []
stopwords = [ "a", "about", "above", "after", "again", "against", "all", "am", "an", "and", "at", "be", "because", "before", "but", "by", "can", "could", "do", "each", "for", "from", "if", "in", "into", "is", "it", "of", "on", "or", "over", "so", "that", "the", "their", "they", "this", "to", "too", "us", "was", "were", "with", "without", "would", "you"]
print(len(stopwords))
```

```
# Expected Output
```

```
# 153
```

```
↳ 153
```

```
with open("/tmp/bbc-text.csv", 'r') as csvfile:
    reader = csv.reader(csvfile, delimiter=',')
    next(reader)
    for label, sentence in reader:
        labels.append(label)
        sentence_to_remove_stop_words = sentence
        for stop_word in stopwords:
            word_token = " " + stop_word + " "
            sentence_to_remove_stop_words = sentence_to_remove_stop_words.replace(word_token, " ")
        sentences.append(sentence_to_remove_stop_words)
```

```
print(len(labels))
```

```
print(len(sentences))
```

```
print(sentences[0])
```

```
# Expected Output
```

```
# 2225
```

```
# 2225
```

```
# tv future hands viewers home theatre systems plasma high-definition tvs digital video rec
```

```
↳ 2225
```

```
2225
```

```
tv future hands viewers home theatre systems plasma high-definition tvs digital video
```

```
#truncate decimals
```

```
train_size = int(len(sentences) * training_portion)
```

```
train_sentences = sentences[:train_size]
```

```
train_labels = labels[:train_size]
```

```
validation_sentences = sentences[train_size:]
```

```
validation_labels = labels[train_size:]
```

```
print(train_size)
```

```
print(len(train_sentences))
```

```
print(len(train_labels))
```

```
print(len(validation_sentences))
```

```
print(len(validation_labels))
```

```
# Expected output (if training_portion=.8)
```

```
# 1780
```

```
# 1780
```

```
# 1780
```

```
# 445
```

```
# 445
```

```

1780
1780
1780
445
445

```

```

tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)
tokenizer.fit_on_texts(train_sentences)
word_index = tokenizer.word_index

```

```

train_sequences = tokenizer.texts_to_sequences(train_sentences)
train_padded = pad_sequences(train_sequences, padding=padding_type, maxlen=max_length, truncat

```

```

print(len(train_sequences[0]))
print(len(train_padded[0]))

```

```

print(len(train_sequences[1]))
print(len(train_padded[1]))

```

```

print(len(train_sequences[10]))
print(len(train_padded[10]))

```

```

# Expected Ouput
# 449
# 120
# 200
# 120
# 192
# 120

```

```

449
120
200
120
192
120

```

```

validation_sequences = tokenizer.texts_to_sequences(validation_sentences)
validation_padded = pad_sequences(validation_sequences, padding=padding_type, maxlen=max_leng

```

```

print(len(validation_sequences))
print(validation_padded.shape)

```

```

# Expected output
# 445
# (445, 120)

```

```

445
(445, 120)

```

```

label_tokenizer = Tokenizer()
label_tokenizer.fit_on_texts(labels)

training_label_seq = np.array(label_tokenizer.texts_to_sequences(train_labels))
validation_label_seq = np.array(label_tokenizer.texts_to_sequences(validation_labels))

print(training_label_seq[0])
print(training_label_seq[1])
print(training_label_seq[2])
print(training_label_seq.shape)

print(validation_label_seq[0])
print(validation_label_seq[1])
print(validation_label_seq[2])
print(validation_label_seq.shape)

# Expected output
# [4]
# [2]
# [1]
# (1780, 1)
# [5]
# [4]
# [3]
# (445, 1)

```

```

[4]
[2]
[1]
(1780, 1)
[5]
[4]
[3]
(445, 1)

```

```

model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(24, activation='relu'),
    tf.keras.layers.Dense(6, activation='softmax')
])
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.summary()

```

```

# Expected Output
# Layer (type)                Output Shape          Param #
# =====
# embedding (Embedding)       (None, 120, 16)      16000
#
# global_average_pooling1d (Gl (None, 16)              0
#
# dense (Dense)               (None, 24)           408

```

```
#
# dense_1 (Dense)          (None, 6)          150
# =====
# Total params: 16,558
# Trainable params: 16,558
# Non-trainable params: 0
```

📄 Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 120, 16)	160000

global_average_pooling1d (Gl	(None, 16)	0

dense (Dense)	(None, 24)	408

dense_1 (Dense)	(None, 6)	150
=====		
Total params: 160,558		
Trainable params: 160,558		
Non-trainable params: 0		

```
num_epochs = 30
history = model.fit(train_padded, training_label_seq, epochs=num_epochs, validation_data=(val
```

📄

```

Epoch 1/30
56/56 [=====] - 1s 9ms/step - loss: 1.7548 - accuracy: 0.2551 -
Epoch 2/30
56/56 [=====] - 0s 5ms/step - loss: 1.6230 - accuracy: 0.3635 -
Epoch 3/30
56/56 [=====] - 0s 5ms/step - loss: 1.3826 - accuracy: 0.5303 -
Epoch 4/30
56/56 [=====] - 0s 5ms/step - loss: 1.0676 - accuracy: 0.8051 -
Epoch 5/30
56/56 [=====] - 0s 5ms/step - loss: 0.7557 - accuracy: 0.9360 -
Epoch 6/30
56/56 [=====] - 0s 4ms/step - loss: 0.4980 - accuracy: 0.9753 -
Epoch 7/30
56/56 [=====] - 0s 5ms/step - loss: 0.3178 - accuracy: 0.9882 -
Epoch 8/30
56/56 [=====] - 0s 5ms/step - loss: 0.2067 - accuracy: 0.9927 -
Epoch 9/30
56/56 [=====] - 0s 5ms/step - loss: 0.1407 - accuracy: 0.9949 -
Epoch 10/30
56/56 [=====] - 0s 4ms/step - loss: 0.1000 - accuracy: 0.9972 -
Epoch 11/30
56/56 [=====] - 0s 5ms/step - loss: 0.0740 - accuracy: 0.9983 -
Epoch 12/30
56/56 [=====] - 0s 5ms/step - loss: 0.0562 - accuracy: 0.9989 -
Epoch 13/30
56/56 [=====] - 0s 5ms/step - loss: 0.0438 - accuracy: 1.0000 -
Epoch 14/30
56/56 [=====] - 0s 5ms/step - loss: 0.0348 - accuracy: 1.0000 -
Epoch 15/30
56/56 [=====] - 0s 6ms/step - loss: 0.0282 - accuracy: 1.0000 -
Epoch 16/30
56/56 [=====] - 0s 5ms/step - loss: 0.0232 - accuracy: 1.0000 -
Epoch 17/30
56/56 [=====] - 0s 5ms/step - loss: 0.0192 - accuracy: 1.0000 -

```

```
import matplotlib.pyplot as plt
```

```

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_'+string])
    plt.show()


```

```

plot_graphs(history, "accuracy")
plot_graphs(history, "loss")

```





```
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
```

```
def decode_sentence(text):
    return ' '.join([reverse_word_index.get(i, '?') for i in text])
```

Epochs

```
e = model.layers[0]
weights = e.get_weights()[0]
print(weights.shape) # shape: (vocab_size, embedding_dim)
```

```
# Expected output
# (1000, 16)
```

```
↳ (10000, 16)
```

```
import io
```

```
out_v = io.open('vecs.tsv', 'w', encoding='utf-8')
out_m = io.open('meta.tsv', 'w', encoding='utf-8')
for word_num in range(1, vocab_size):
    word = reverse_word_index[word_num]
    embeddings = weights[word_num]
    out_m.write(word + "\n")
    out_v.write('\t'.join([str(x) for x in embeddings]) + "\n")
out_v.close()
out_m.close()
```

```
try:
```

```
from google.colab import files
except ImportError:
    pass
else:
    files.download('vecs.tsv')
    files.download('meta.tsv')
```

