

```

#@title Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
#
# https://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.

```



```

import tensorflow as tf

from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.layers import Embedding, LSTM, Dense, Bidirectional
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
import numpy as np

!wget --no-check-certificate \
  https://storage.googleapis.com/laurencemoroney-blog.appspot.com/irish-lyrics-eof.txt \
  -O /tmp/irish-lyrics-eof.txt

tokenizer = Tokenizer()

data = open('/tmp/irish-lyrics-eof.txt').read()

corpus = data.lower().split("\n")

tokenizer.fit_on_texts(corpus)
total_words = len(tokenizer.word_index) + 1

```

```
print(tokenizer.word_index)
print(total_words)
```

```
{ 'the': 1, 'and': 2, 'i': 3, 'to': 4, 'a': 5, 'of': 6, 'my': 7, 'in': 8, 'me': 9, 'for': 10 }
```

```
input_sequences = []
for line in corpus:
    token_list = tokenizer.texts_to_sequences([line])[0]
    for i in range(1, len(token_list)):
        n_gram_sequence = token_list[:i+1]
        input_sequences.append(n_gram_sequence)

# pad sequences
max_sequence_len = max([len(x) for x in input_sequences])
input_sequences = np.array(pad_sequences(input_sequences, maxlen=max_sequence_len, padding='p

# create predictors and label
xs, labels = input_sequences[:, :-1], input_sequences[:, -1]

ys = tf.keras.utils.to_categorical(labels, num_classes=total_words)
```

```
print(tokenizer.word_index['in'])
print(tokenizer.word_index['the'])
print(tokenizer.word_index['town'])
print(tokenizer.word_index['of'])
print(tokenizer.word_index['athy'])
```

Saving... ✕

```
print(tokenizer.word_index['lanigan'])
```

8
1
71
6
713
39
1790
1791

```
print(xs[6])
```

⇒ $[0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 2]$

```
print(ys[6])
```

→ [0. 0. 0. ... 0. 0. 0.]

```
print(xs[5])
print(ys[5])
```

```
↳ [ 0 0 0 0 0 0 0 0 0 0 51 12 96 1217 48
    2]
    [0. 0. 0. ... 0. 0. 0.]
```

```
print(tokenizer.word_index)
```

```
↳ {'the': 1, 'and': 2, 'i': 3, 'to': 4, 'a': 5, 'of': 6, 'my': 7, 'in': 8, 'me': 9, 'for'
```

```
model = Sequential()
model.add(Embedding(total_words, 100, input_length=max_sequence_len-1))
model.add(Bidirectional(LSTM(150)))
model.add(Dense(total_words, activation='softmax'))
adam = Adam(lr=0.01)
model.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
#earlystop = EarlyStopping(monitor='val_loss', min_delta=0, patience=5, verbose=0, mode='auto')
history = model.fit(xs, ys, epochs=100, verbose=1)
#print model.summary()
print(model)
```

```
↳
```

Saving...



```
377/377 [=====] - 4s 9ms/step - loss: 0.8623 - accuracy: 0.7695
Epoch 73/100
377/377 [=====] - 4s 9ms/step - loss: 0.9539 - accuracy: 0.7515
Epoch 74/100
377/377 [=====] - 4s 9ms/step - loss: 0.9131 - accuracy: 0.7565
Epoch 75/100
377/377 [=====] - 4s 9ms/step - loss: 0.8653 - accuracy: 0.7691
Epoch 76/100
377/377 [=====] - 3s 9ms/step - loss: 0.8391 - accuracy: 0.7745
Epoch 77/100
377/377 [=====] - 3s 9ms/step - loss: 0.8353 - accuracy: 0.7775
Epoch 78/100
377/377 [=====] - 4s 9ms/step - loss: 0.7672 - accuracy: 0.7937
Epoch 79/100
377/377 [=====] - 3s 9ms/step - loss: 0.7730 - accuracy: 0.7944
Epoch 80/100
377/377 [=====] - 4s 9ms/step - loss: 0.7834 - accuracy: 0.7927
Epoch 81/100
377/377 [=====] - 4s 10ms/step - loss: 0.8452 - accuracy: 0.7835
Epoch 82/100
377/377 [=====] - 4s 9ms/step - loss: 0.8641 - accuracy: 0.7738
Epoch 83/100
377/377 [=====] - 4s 9ms/step - loss: 0.8770 - accuracy: 0.7682
Epoch 84/100
377/377 [=====] - 4s 10ms/step - loss: 0.8845 - accuracy: 0.7635
Epoch 85/100
377/377 [=====] - 4s 10ms/step - loss: 0.8562 - accuracy: 0.7695
Epoch 86/100
377/377 [=====] - 4s 10ms/step - loss: 0.8835 - accuracy: 0.7695
Epoch 87/100
377/377 [=====] - 4s 10ms/step - loss: 0.8686 - accuracy: 0.7745
Epoch 88/100
377/377 [=====] - 4s 9ms/step - loss: 0.8464 - accuracy: 0.7796
Epoch 89/100
377/377 [=====] - 4s 9ms/step - loss: 0.8206 - accuracy: 0.7847
Epoch 90/100
377/377 [=====] - 4s 9ms/step - loss: 0.7915 - accuracy: 0.7927
Epoch 91/100
377/377 [=====] - 4s 9ms/step - loss: 0.8604 - accuracy: 0.7805
Epoch 92/100
377/377 [=====] - 4s 9ms/step - loss: 0.8733 - accuracy: 0.7727
Epoch 93/100
377/377 [=====] - 3s 9ms/step - loss: 0.8570 - accuracy: 0.7736
Epoch 94/100
377/377 [=====] - 4s 9ms/step - loss: 0.8545 - accuracy: 0.7745
Epoch 95/100
377/377 [=====] - 4s 9ms/step - loss: 0.8546 - accuracy: 0.7745
Epoch 96/100
377/377 [=====] - 4s 9ms/step - loss: 0.8316 - accuracy: 0.7835
Epoch 97/100
377/377 [=====] - 4s 9ms/step - loss: 0.8193 - accuracy: 0.7832
Epoch 98/100
377/377 [=====] - 4s 9ms/step - loss: 0.8595 - accuracy: 0.7788
Epoch 99/100
377/377 [=====] - 4s 9ms/step - loss: 0.8275 - accuracy: 0.7801
Epoch 100/100
377/377 [=====] - 4s 9ms/step - loss: 0.8011 - accuracy: 0.7882
tf.nn.conv2d: Out of memory: Requested 128000000 bytes for the operation. This operation needs in total 128000000 bytes of memory. Please consider reducing the batch size or the number of filters. Sequential object at 0x7f86f2010200
```

Saving...



```
<tensorflow.python.keras.engine.sequential.Sequential object at 0x7f8b72010208>
```

Saving...

